



NTNU

Innovation and Creativity

***Offload Compiler Runtime for the Intel® Xeon
Phi™ Coprocessor***

C.J. Newton, R. Deodhar, S. Dmitriev, et al.

TDT24 Paper presentation

Christian Chavez

Institutt for Datateknikk og Informasjonsvitenskap

November 10, 2014

Outline

Introduction

Background: GPGPUs

Background: Xeon Phi

Runtime Offloading: Tools

Runtime Offloading: Application tests

Runtime Offloading: Test results


Conclusion

Motivation

What happens when you have a codebase containing millions of lines of code?

Motivation

What happens when you have a codebase containing millions of lines of code?



NTNU

Innovation and Creativity

GPGPUs

So, how can *GPGPUs* help?

1. Split up and run sub-parts of the compilation **concurrently** with host (CPU)

GPGPUs

So, how can *GPGPUs* help?

1. Split up and run sub-parts of the compilation **concurrently** with host (CPU)
2. Run **computationally heavy** parts of the compiler algorithms, while the CPU works on the more serial ones.

Intel coprocessors

Intel wants to enter this market with the help of the **Xeon Phi** processor family

Intel coprocessors

Intel wants to enter this market with the help of the **Xeon Phi** processor family

- Intended as a **competitor** against Nvidia CUDA GPGPUs in the HPC market

Xeon Phi drawbacks

Obvious *drawbacks* with Xeon Phi (compared with the CUDA GPGPUs)

1. Not as **mature** ecosystem/tools/environment

Xeon Phi drawbacks

Obvious *drawbacks* with Xeon Phi (compared with the CUDA GPGPUs)

1. Not as **mature** ecosystem/tools/environment
2. Still requires a high degree of **computation to communication** to be beneficial to run on the Phi instead of the CPU

Xeon Phi drawbacks

Obvious *drawbacks* with Xeon Phi (compared with the CUDA GPGPUs)

1. Not as **mature** ecosystem/tools/environment
2. Still requires a high degree of **computation to communication** to be beneficial to run on the Phi instead of the CPU
3. **Still new and relatively untested/unused wrt. compilers**

Xeon Phi drawbacks

Obvious *drawbacks* with Xeon Phi (compared with the CUDA GPGPUs)

1. Not as **mature** ecosystem/tools/environment
2. Still requires a high degree of **computation to communication** to be beneficial to run on the Phi instead of the CPU
3. Still new and relatively untested/unused wrt. compilers
4. **Being Intel, everything is “canned” (not open source)**

Xeon Phi benefits

Obvious *benefits* with Xeon Phi (compared with the CUDA GPGPUs)

1. The architecture supports regular Intel CPU instructions¹

¹Not confirmed exactly how much

Xeon Phi benefits

Obvious *benefits* with Xeon Phi (compared with the CUDA GPGPUs)

1. The architecture supports regular Intel CPU instructions¹
 - The Xeon Phi does not require its own codebase for being run.

¹Not confirmed exactly how much

Runtime offloading tools

The papers focus is on the following tools:

- Intel Manycore Platform Software Stack (MPSS)

Runtime offloading tools

The papers focus is on the following tools:

- Intel Manycore Platform Software Stack (MPSS)
 - Intel Coprocessor Offload Infrastructure (COI)

Runtime offloading tools

The papers focus is on the following tools:

- Intel Manycore Platform Software Stack (MPSS)
 - Intel Coprocessor Offload Infrastructure (COI)
- COI utilizes:

Runtime offloading tools

The papers focus is on the following tools:

- Intel Manycore Platform Software Stack (MPSS)
 - Intel Coprocessor Offload Infrastructure (COI)
- COI utilizes:
 - **Library calls**

Runtime offloading tools

The papers focus is on the following tools:

- Intel Manycore Platform Software Stack (MPSS)
 - Intel Coprocessor Offload Infrastructure (COI)
- COI utilizes:
 - Library calls
 - **Language pragmas**

Runtime offloading tools

The papers focus is on the following tools:

- Intel Manycore Platform Software Stack (MPSS)
 - Intel Coprocessor Offload Infrastructure (COI)
- COI utilizes:
 - Library calls
 - Language pragmas
 - **Language keywords**

Runtime offloading tools

The papers focus is on the following tools:

- Intel Manycore Platform Software Stack (MPSS)
 - Intel Coprocessor Offload Infrastructure (COI)
- COI utilizes:
 - Library calls
 - Language pragmas
 - Language keywords
 - A FIFO command queue (dataobject: COIPipeline)



NTNU

Innovation and Creativity

Runtime offloading tools

The papers focus is on the following tools:

- Intel Manycore Platform Software Stack (MPSS)
 - Intel Coprocessor Offload Infrastructure (COI)
- COI utilizes:
 - Library calls
 - Language pragmas
 - Language keywords
 - A FIFO command queue (dataobject: COIPipeline)
- OpenMP

Platform specifications

Test platform specifications:

Table 1: Platform configuration parameters	
Host	SNB-EP (2 sockets) 2.6 GHz, Intel® Xeon® E5-2670, Crown Pass Platform
Copro-processor	Pre-production Intel® Xeon Phi™ coprocessor, 61 4-thread cores, 1.09GHz, 5.5GTransfers/s, 8GB
Host OS	RHEL 6.2, kernel 2.6.32-220.el6.x86_64
Compiler	Composer XE Beta
MPSS	2.1.3653-8

Two sets of tests:

1. A test spanning a variety of application domains, coming from customers.

Two sets of tests:

1. A test spanning a variety of application domains, coming from customers.
 - Hence, some names are occluded in deference to customers.

Two sets of tests:

1. A test spanning a variety of application domains, coming from customers.
 - Hence, some names are occluded in deference to customers.
 - **The tests are meaningfully representative to said customers.**

Two sets of tests:

1. A test spanning a variety of application domains, coming from customers.
 - Hence, some names are occluded in deference to customers.
 - The tests are meaningfully representative to said customers.
2. **An array of benchmarks from Scalable Heterogeneous Computing (SHOC) benchmarks suite.**



NTNU

Innovation and Creativity

Two sets of tests:

1. A test spanning a variety of application domains, coming from customers.
 - Hence, some names are occluded in deference to customers.
 - The tests are meaningfully representative to said customers.
2. An array of benchmarks from Scalable Heterogeneous Computing (*SHOC*) benchmarks suite.
 - Only used tests that were not exclusive to OpenCL and CUDA



NTNU

Innovation and Creativity

Two sets of tests:

1. A test spanning a variety of application domains, coming from customers.
 - Hence, some names are occluded in deference to customers.
 - The tests are meaningfully representative to said customers.
2. An array of benchmarks from Scalable Heterogeneous Computing (*SHOC*) benchmarks suite.
 - Only used tests that were not exclusive to OpenCL and CUDA
 - **Only used tests that have both offload and native versions**



NTNU

Innovation and Creativity

Two sets of tests:

1. A test spanning a variety of application domains, coming from customers.
 - Hence, some names are occluded in deference to customers.
 - The tests are meaningfully representative to said customers.
2. An array of benchmarks from Scalable Heterogeneous Computing (*SHOC*) benchmarks suite.
 - Only used tests that were not exclusive to OpenCL and CUDA
 - Only used tests that have both offload and native versions
 - Results were measured by Intel in August 2012.



NTNU

Innovation and Creativity

Two sets of tests:

1. A test spanning a variety of application domains, coming from customers.
 - Hence, some names are occluded in deference to customers.
 - The tests are meaningfully representative to said customers.
2. An array of benchmarks from Scalable Heterogeneous Computing (*SHOC*) benchmarks suite.
 - Only used tests that were not exclusive to OpenCL and CUDA
 - Only used tests that have both offload and native versions
 - Results were measured by Intel in August 2012.
 - The benchmark application *Triad* was also added, even though it has no native version

Customer tests results

Table 2: Offload performance and overheads for workloads

Workload [some names occluded, pending customer approval]	3DFD PDE Stencil	Convolution Resampling	Hogbom-Clean	QR	Iterative closest point, DP	Adaptive Sparse Grid	Black Scholes Compute SP
Domain	Seismic	Astronomy	Astronomy	Physics	Manufacturing	Physics	Financial
Speedup with offload	2.03	1.56	2.31	1.40	1.54	1.32	6.92
Compute % of total execution time	97.7	44.3	72.9	97.3	94.5	95.3	99.4
Host offload overhead factor, with (top) & without (bottom) initialization	0.21	1.77	0.44	0.03	0.02	0.06	0.00
	0.18	1.60	0.25	0.01	0.01	0.05	0.00
Computation/Communication ratio	5.42	0.62	3.99	68.6	90.1	19.7	1640



NTNU

Innovation and Creativity

SHOC tests results

Table 3: Offload performance and overheads and native performance for SHOC													
Workload	FFT-DP	FFT-SP	GEMM-DP	GEMM-SP	MD-DP	MD-SP	Reduction-DP	Reduction-SP	S3D-DP	S3D-SP	SCAN	Sort	SPMV-DP
Data set	N= 1677216	N= 3354432	2048 x 2048	4096 x 4096	73728_atoms	73728_atoms	838608_items	16777216_items	262144_gridPoints	262144_gridPoints	8388608_items	62451_rows	4007383_elements
Speedup with offload (top) and native only	0.30	0.45	3.30	3.25	0.81	0.44	0.20	0.20	0.95	1.03	0.25	0.21	0.07
% execution time in offload	3.58	3.62	33.0	59.1	19.3	0.26	37.3	19.0	7.53	6.92	68.0	17.4	8.22
Host offload overhead, with & without init	13.9	14.8	1.08	0.34	2.42	17.1	0.92	2.42	9.75	9.45	0.19	3.66	6.09
Computation: communication	0.83	0.92	7.74	14.0	1.92	1.85	2.46	2.10	0.20	0.32	56.9	0.40	0.63



NTNU

Innovation and Creativity

Speedups 1/2

- Computation to communication ratio varied from 0.62 to 1640

Speedups 1/2

- Computation to communication ratio varied from 0.62 to 1640
 1. But it was generally over 4, except for the convolution case.

Speedups 1/2

- Computation to communication ratio varied from 0.62 to 1640
 1. But it was generally over 4, except for the convolution case.
 2. The convolution case scaling across threads and SIMD elements on the coprocessor allowed the computation speedup to outweigh the overheads



NTNU

Innovation and Creativity

Speedups 1/2

- Computation to communication ratio varied from 0.62 to 1640
 1. But it was generally over 4, except for the convolution case.
 2. The convolution case scaling across threads and SIMD elements on the coprocessor allowed the computation speedup to outweigh the overheads
- Data movement overhead varied from negligible to 1.77x the computation time

Speedups 1/2

- Computation to communication ratio varied from 0.62 to 1640
 1. But it was generally over 4, except for the convolution case.
 2. The convolution case scaling across threads and SIMD elements on the coprocessor allowed the computation speedup to outweigh the overheads
- Data movement overhead varied from negligible to 1.77x the computation time
- **Black Scholes was the big performance winner, with a 6.92x speedup of the dual-socket SandyBridge**

Speedups 1/2

- Computation to communication ratio varied from 0.62 to 1640
 1. But it was generally over 4, except for the convolution case.
 2. The convolution case scaling across threads and SIMD elements on the coprocessor allowed the computation speedup to outweigh the overheads
- Data movement overhead varied from negligible to 1.77x the computation time
- Black Scholes was the big performance winner, with a 6.92x speedup of the dual-socket SandyBridge
- SHOC shows no correlation between speedup and computation to communication.

Speedups 1/2

- Computation to communication ratio varied from 0.62 to 1640
 1. But it was generally over 4, except for the convolution case.
 2. The convolution case scaling across threads and SIMD elements on the coprocessor allowed the computation speedup to outweigh the overheads
- Data movement overhead varied from negligible to 1.77x the computation time
- Black Scholes was the big performance winner, with a 6.92x speedup of the dual-socket SandyBridge
- SHOC shows no correlation between speedup and computation to communication.
 - Paper lists this as potential future work.
“Broader investigation is needed on that”



NTNU

Innovation and Creativity

Speedups 2/2

“No single factor determines offload performance”

- Data movement overhead varied from negligible to 1.77x the computation time

Offload speedups are most correlated with the following:

Speedups 2/2

“No single factor determines offload performance”

- Data movement overhead varied from negligible to 1.77x the computation time

Offload speedups are most correlated with the following:

1. **Inverse of offload overheads**

Speedups 2/2

“No single factor determines offload performance”

- Data movement overhead varied from negligible to 1.77x the computation time

Offload speedups are most correlated with the following:

1. Inverse of offload overheads
2. % of execution time in offload (which includes coprocessor invocation and data movement)

Speedups 2/2

“No single factor determines offload performance”

- Data movement overhead varied from negligible to 1.77x the computation time

Offload speedups are most correlated with the following:

1. Inverse of offload overheads
2. % of execution time in offload (which includes coprocessor invocation and data movement)
3. **Native performance (when there is no execution on the host at all)**



NTNU

Innovation and Creativity

Conclusion

- The Xeon Phi in conjunction with the platform used in the paper showed benchmark application speedups from 1.3x to 6.9 on “customer-relevant” examples spanning different application domains
- Offload is not always profitable, as SHOC showed
- Inhibitions or enhancements of speedup this paper has explored:
 1. When response time is of concern, there must be a speedup from native execution on only the coprocessor, relative to the execution on the host
 2. Ratio of computation to communication must be generally high
 3. Offload runtime overheads must be small relative to computation



NTNU

Innovation and Creativity

Sources

- *Offload Compiler Runtime for the Intel® Xeon Phi™ Coprocessor*, Newton, R. Deodhar, S. Dmitriev, et al.²
- Randall Munroe, xkcd³

²<https://software.intel.com/sites/default/files/article/366893/offload-runtime-for-the-intelr-xeon-phitm-coprocessor.pdf>

³<http://xkcd.com/303/>



NTNU

Innovation and Creativity